

# HOW TO STAY ALIVE EVEN WHEN OTHERS GO DOWN

WRITING AND TESTING  
RESILIENT APPLICATIONS





***shopify***



... even when others go down.

Who is “others”?

“We don’t have a service-oriented architecture.  
Just a monolithic Rails application.”



# "SERVICES" YOU MIGHT HAVE

- Database
- Cache
- Sessions and authentication
- Asset file storage
- Search
- Background job queue
- Analytics and tracking
- Image processing
- Recommendations
- ...

# RESILIENCY PATTERNS


- Timeouts
- Circuit Breakers
- Bulkheading
- Fallbacks
- Testing
- Chaos Monkey

# CAPACITY







# NGINX+UNICORN ARCHITECTURE



# CAPACITY



# TIMEOUTS: FAIL FAST

Failing can mean different things.

Connection refused after 0.1s?

Connection timed out after 5s?

Connection established but read timed out after 30s?

Server returns unparseable data?

# TIMEOUTS: FAIL FAST

Ruby gem	Default timeout	Better timeout
Unicorn	60s	~5s
Net::HTTP	60s	~2s
mysql2	none	~1s
redis-rb	none	~0.5s
AWS::S3	60s	~2s
memcached	0.5s	
...		



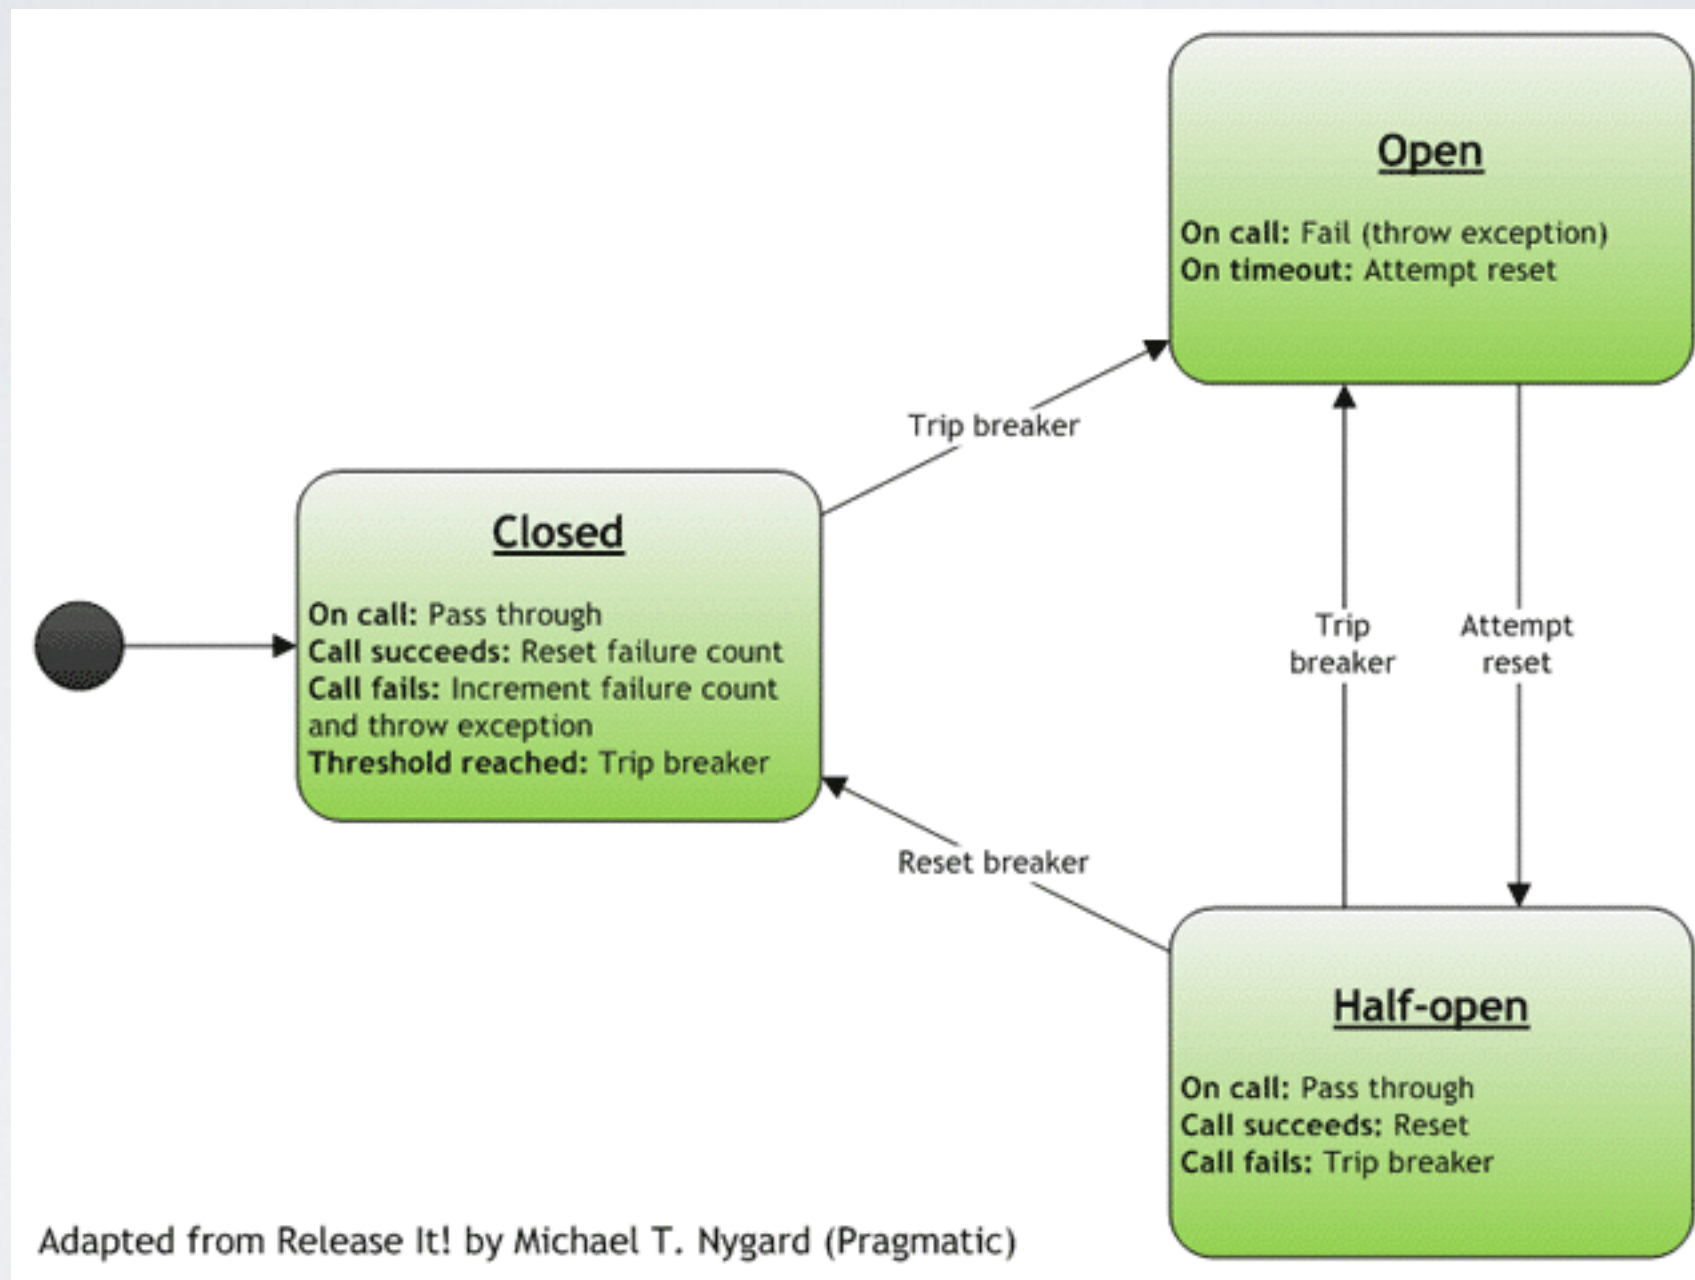
# REAL-WORLD EXAMPLES

- Redis KEYS command
- Slow MySQL queries (missing indexes, expensive JOINS, etc.)
- “Stop-the-world” garbage collection
- Heroku app runs out of dynos
- Datacenter tech shutting off the wrong machine
- Hadoop node on same switch
- Network packet loss, retransmits, etc.

# CIRCUIT BREAKERS

- Observation: If the operation **failed**, retrying it immediately will likely **fail again**.
- Idea: Keep track of errors and **stop trying for a while** if threshold is reached.
- Each service (each backend) has its own “circuit”.

# CIRCUIT BREAKERS



# IMPLEMENTATION TIPS

- Implement at the driver level, not in the consumer.
- Different code paths can **share the same circuit** if they are hitting the same service (or even different services with same backend).
- Treat “**CircuitOpen**” exceptions the same way you treat other backend errors (fallbacks).



# FALLBACKS: FAIL GRACEFULLY

- “Don't be defensive.”
- “Don't need to rescue this exception. If Redis is down, we have bigger problems.”
- If a service is unavailable for whatever reason, try to return a reasonable fallback value.

## SUNGLASSES

CartService (redis1)

A-Z Z-A \$-\$\$ \$\$-\$



SessionService  
(redis2)

AssetService (AWS::S3)



GALAXY PURPLE PEEPS  
**\$199.00 AUD**

Sold Out - Not Available :(

InventoryService (redis3)



Fragment cache  
(memcached)

KOI PEEPS - LIMITED  
**\$199.00 AUD**

Available : OSFM

ProductSearchService  
(ElasticSearch)



TARTAN RED PEEPS - LIMITED  
**\$199.00 AUD**

Available : OSFM

# FALLBACKS: FAIL GRACEFULLY

- ElasticSearch: Assume empty result set
- Sessions: Guest checkout
- Personal recommendations: Generic recommendations
- Distributed lock: Assume someone else holds the lock
- A/B testing: Assume control group
- Throttling: Assume unthrottled
- Cache: Assume cache miss
- ...

# IMPLEMENTATION TIPS

- Push fallbacks deep down the stack (can be tricky).
- Write good abstractions (don't rescue Redis errors at the controller level).
- Monitoring and alerting!



```
class SomeDatastoreClient
  class CircuitOpenError < BaseError
    end

    # ...
  end

  class ShoppingCart
    class CartUnavailable < StandardError
      end

      # ...

      def load(cart_id)
        @datastore.get(cart_id)
        rescue SomeDatastoreClient::BaseError
          raise CartUnavailable
        end
      end
    end

    class SomeRailsController

      # ...

      def load_cart
        @cart = ShoppingCart.load(session[:cart_id])
        rescue ShoppingCart::CartUnavailable
          flash[:notice] = 'The cart system is currently unavailable'
          EmptyCart.new
        end
      end
    end
  end
end
```

# BULKHEADING: ISOLATE FAILURES

- Ensure that failures in one component don't cause **cascading failures** in other components.
- Limit concurrent access to shared resources (e.g., using semaphores).
- **Isolate** services from each other.

# BULKHEADING: EXAMPLES

- Concurrency control: Only N workers allowed to talk to a given resource at once, then block.
- Don't share datastore instances between use cases. Use separate processes (or even hardware) instead of logical databases.
- Throttle error reporting jobs.
- MySQL failures shouldn't break pages that don't even use MySQL.

# WHERE TO START?

- Start where the **money** is
- Application hot paths
- Most traffic / most “visible”
- Often forgotten: **Deploys!**
- Don't do it all at once



# RESILIENCY MATRIX

Component / Site area	Storefront	Checkout	Admin
MySQL	OK (if cached)	Down	Down
Redis (Sessions)	Down	OK	Down
Redis (Inventory)	Degraded	Down	OK
AWS S3	OK	OK	Degraded

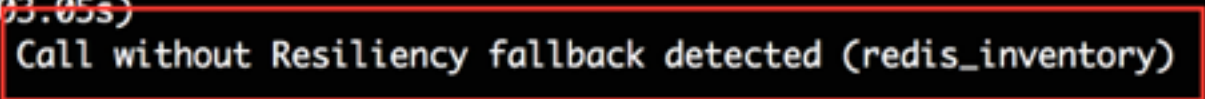
# RESILIENCY TESTING

- Prevent regressions.
- **Also great exploration tool!**
- Use your tests to generate a TODO list for you!
- Then fix iterative, one test class at a time.

# GENERATING A TODO LIST

```
2
3 class CartsControllerTest < ActionController::TestCase
4   include Resiliency::TestWithVerification
5
6
```

```
ERROR["test_create_order_with_non_taxable_line_item", CartsControllerTest, 2015-09-06 20:38:23 +0000]
test_create_order_with_non_taxable_line_item#CartsControllerTest (1441571903.05s)
Resiliency::MissingFallbackError:      Resiliency::MissingFallbackError: Call without Resiliency fallback detected (redis_inventory)
lib/resiliency/verifier.rb:18 in `fallback_verification`
test/support/resiliency/redis_verification.rb:3 in `io`
(redis-rb-977ae2165fce) lib/redis/client.rb:255 in `write`
```



# REGRESSION TESTING

```
test "storefront can serve /products cache hits without mysql" do
  # warm cache
  get '/products'

  toxiproxy(/mysql/).down do
    get '/products'
    assert_response 200
    assert_equal 'hit', response.headers['X-Cache']
  end
end
```


```
test "#get gracefully handles connection error and does not raise" do
  toxiproxy(SessionStore).down do
    session = @session_store.get('my-session-id')
    assert session.empty?
  end
end
```



# CHAOS MONKEYS

- Pull the plug on things in production
- Regularly scheduled “game days”
- Introduce failures (kill -9, iptables, fill up disk, ...)
- Introduce latency, network partitions, ...
- Start with controlled failures  
(only one server, only x% of traffic, ...)
- Once confident: Automate!








# HOMEWORK

- Set timeouts on EVERYTHING.
- Implement the circuit breaker pattern yourself.
- Create a resiliency matrix for your application.
- Verify correctness of matrix by writing regression tests.
- [Bonus] Implement “fail test if service used without fallback” helper.
- [Bonus] Implement concurrency control using semaphores.

# WANT TO LEARN MORE?

- "Release It!"
- Netflix tech blog
- Shopify/toxiproxy
- Shopify/semian
- Shopify tech Blog
- @fw1729 / flo@shopify.com





Thanks for your attention!